Strassen improves transformer computation across all three models

Speedup over Tuned Native | Higher is better | Sequence 1,024 | BF16 output | One TPU v6e

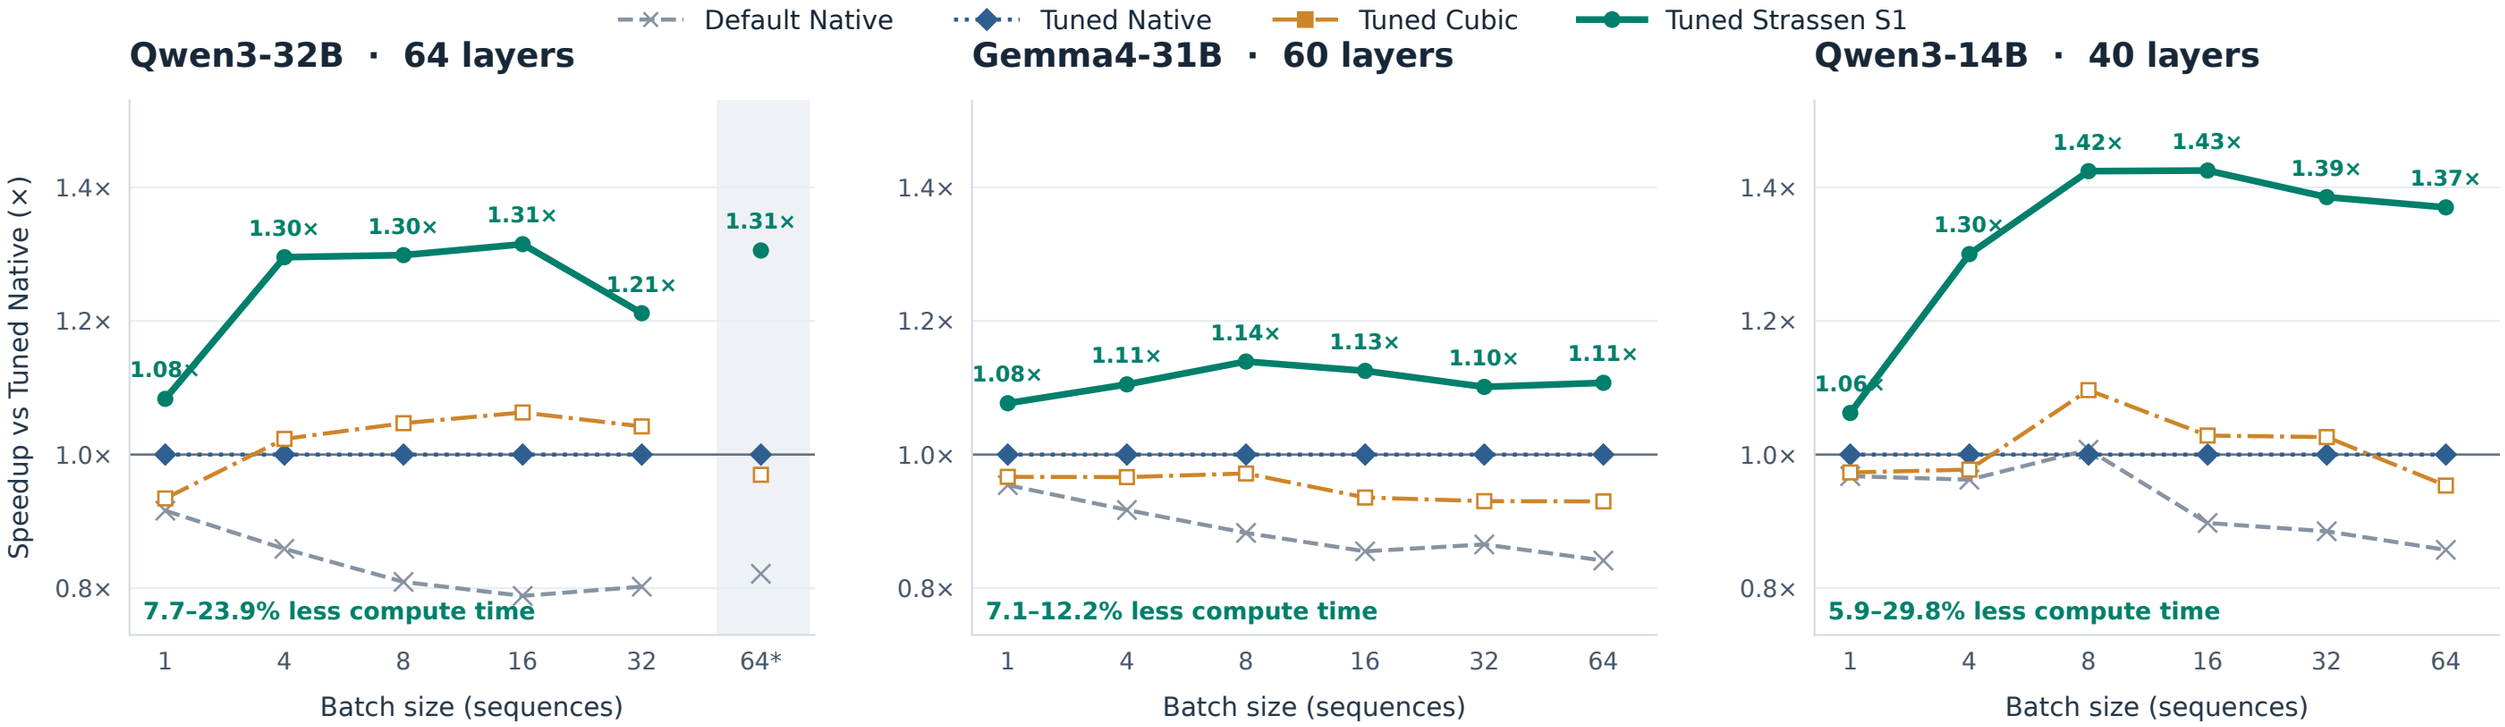

Metric: sum of separately warmed decoder-layer means. Includes layer computation; excludes weight preparation/transfers, embedding/head and compilation.
This is not uninterrupted whole-model latency. Each layer has 2 warmups and 5 measured samples per arm.
* Qwen3-32B B64 uses a newer compiler stack; it is not joined to B32. Within-batch algorithm comparisons are matched.