Perplexity remains close across algorithms

WikiText-2 test | Four contiguous segments per batch | Upper curves nearly overlap; lower panels reveal the differences

Default Native · Tuned Native · Tuned Cubic · Tuned Strassen S1

Qwen3-32B

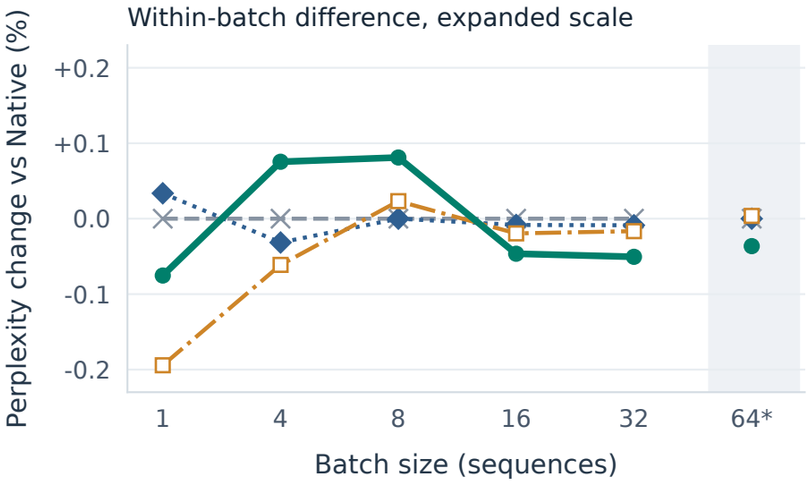

Gemma4-31B

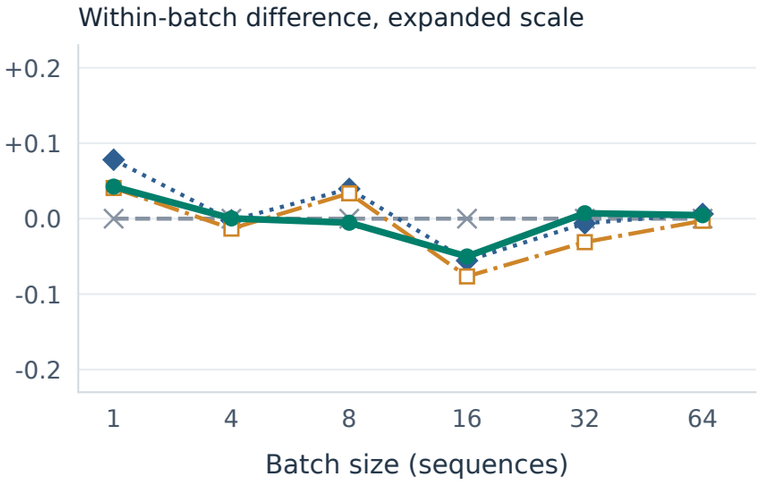

Qwen3-14B

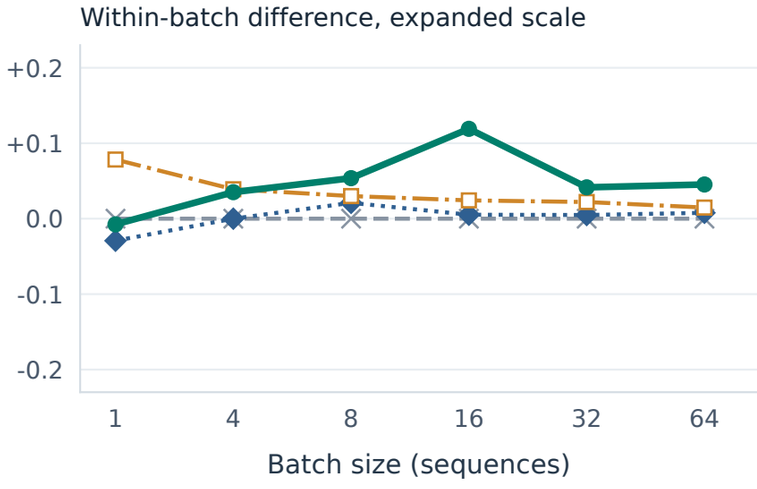

Relative change = 100 × (algorithm perplexity / matched Native perplexity − 1). Positive means higher perplexity; negative means lower.

Corpus spans vary by batch and tokenizers vary by model: compare algorithms within each workload. Close perplexity does not mean identical predictions.

* Qwen3-32B B64 uses a newer compiler stack; it is not joined to B32. Within-batch algorithm comparisons are matched.